



Facts Data Modelling - Day 1

# class	4
☰ Topic	Data Modelling
📌 Notes Status	completed
📅 Completion Date	@December 2, 2024

What is a Fact?

- Something that happened or occurred.
- It's **atomic**, cannot be broken down into a smaller piece
- Facts are like past, **they don't change**

Example:

- User logs in
- Transaction is made

What makes Fact modelling hard?

- Huge volume of data and keeps scaling. Generally 10-100X of dimension data.
- How much fact data is gonna go up depends on how many actions a dimension can make (Ex: Facebook, 25 notifications per user, 2B active users daily = 50B notifications/ day).
- You need a context for effective analysis

- Notification sent alone doesn't mean anything, Notification sent → Clicked → Purchased makes more sense. Conversion funnels .
- CTR alone doesn't make much sense. CTR along with region, CTR is more in US than Canada(ex). Something to do analysis on.



Google tried 40 different shades of blue on the links on their website and one shade of blue had more conversion rates resulting in some **millions of dollars increment in revenue.**

- Duplicates are way more common in fact data than in dimensional data.
 - Duplicate due to some logging bug.
 - Genuine duplicates. Ex: Clicking a notification now and then clicking again 1 hour later. This will make CTR look inflated and results in wrong interpretation of the data.

How Facts Modelling work?

Both Normalization and Denormalization have a place in Facts Data Modelling.

Facts and raw log are not the same thing

Raw logs

- Owned by engineers that don't have strong data skills.
- Ugly schemas designed for online systems.
- Potentially contains duplicates and quality errors.
- Shorter retention

Fact Data

- Nice and descriptive column names.
- Quality guarantees like uniqueness, completeness, etc.
- Longer retention.

Trust in the Fact Data should be in order of magnitude, higher than the trust in raw logs.



One of the best thing you can do as a Data Engineer is to convert those raw logs into a highly trusted fact data

Understanding Who, What, Where, When and How of Facts

Who

- ID, who is the part of the event.

Where

- Location
- Where in the app, PC, Tablet, Phone.

How

- How and Where are sometimes similar, app, PC, Tablet, Phone.

What

- What fields are fundamentally part of the nature of the fact.
- In notification world : Sent, Clicked, Delivered.

When

- Mostly event_timestamp or event_date.



Important Note: Timestamp should be common for all the data, preferably UTC

- Fact tables should have some quality guarantees. The What and When field should never be NULL, otherwise the data cannot be.
- Fact data should be smaller than raw logs.

- Fact data should parse out hard to understand columns. Should be STR, INT, DOUBLE, etc.

"Your job as a Data Engineer is to make the data fun and delightful to use. If you are not doing that, you're probably a bad data engineer." 😊

When should you model in Dimensions (When to de-normalize)

- **[IMPORTANT]** Network Logs pipeline example from Netflix (Embed the URL)
- Broadcast join is optimal if the smaller table is less than 5-6 GB.

"A lot of time the impact of you have as a Data Engineer is not writing and optimizing a pipeline. It's actually going upstream and solving the problem at the source."

How does logging fit into Fact data?

- Brings critical context for your fact data. Usually done in collab with system engineers.
- Do not log everything. If it's just in case, it will never be used.
- **[IMPORTANT]** Logging should conform to some sort of schema or contract.
- Thrift Schema: Specification that is language-agnostic. ? (Need more info).

Potential options when working with high volume fact data

- Do we need to use all the data or can we filter some of it?
- Sampling: Does not work for all use cases, works best for metric-driven use-cases.
- **Bucketing**
 - Fact data can be bucketed on **WHO** i.e. ID or a unique identifier.

- Bucket joins can be much faster as they minimize shuffle.
- **Sorted merge bucket (SMB) joins** can do joins **without shuffle** at all. This was more used during Hive days, in Spark shuffling is more optimized. But it's not gone, **it is still relevant**.

Fact Data Retention

- Any fact tables > 100TB had ~14 days or less retention.
- If size < 10 TB, retention does not matter as long as PII are masked.
- Depends on the company policy.

Deduplication of Fact Data


- Think about the timeframe.
- Window of deduplication : hour, day, week? how to decide?

Intraday deduping options

1. Streaming

- Captures most duplicates in a very effective manner for a short time
- A large memory of duplicate events happen within the short interval of the first event (from notification example).
- Entire duplicates can be harder for streaming because it needs to hold on to such a huge data. (This depends on the use case. If you have the smaller data, we can hold more).

2. Microbatch

Example below. 

Hourly Microbatch Dedupe (Zach's solution at Facebook)

- Reduced the latency from 9 hours to 1 hour ⚡

Implementation

Step 1

- Collect each hour data and dedupe with GROUP BY.

- Use SUM and COUNT to aggregate duplicates, use COLLECT_LIST to collect metadata about the duplicates that might be different.

Step 2

- Dedupe between hours with FULL OUTER JOIN
 - Use left.value + right.value to keep duplicates aggregation correctly counting or CONCAT to build a continuous list (What ?)

Step 3

- Keep merging until we have a final merged table.

